

EE 5450 Project 02: Convolution Neural Network Classification of the Animals Data Set

David R. Mohler

April 19, 2018

1 Introduction

The rebirth of Neural Networks (NNs) and Machine Learning (ML) has offered revolutionary approaches to intelligent tasks previously unattainable in the field of image processing and computer vision. A major area of research continues to be the automated classification of images. Image classification through classical means of image processing tend to suffer due to problems with the wide variance in images to be classified and how their subject matter may be transformed from image to image. The discovery of Convolution Neural Networks (CNNs) lead to a massive increase in the accuracy of which images could be classified in an automated fashion.

For this project we will demonstrate the development of several convolution neural networks that vary across a wide number of parameters for the purpose of classifying color images of three different types of animals. The CNN is tasked with classifying images of cats, dogs, and pandas, typical examples of images in the data set can be seen below. Through the variation of network parameters such as the number of layers and the amount of neurons within them, activation functions, learning rates, optimization routines, data augmentation, and other hyper parameters, we display the ability to increase the classification performance of the network.



2 Methods and Results

After establishing an initial understanding of the basic **ShallowNet** animal classification network that was provided, we begin with a relatively simple and small expansion to include fundamental elements of CNNs, once a working network is establish we are able to tune the parameters and observe the effects of varying the optimization techniques. Once a candidate optimizer has been selected we proceed to modify and tune the associated hyperparameters such as the learning rate (when applicable), batch size, number of training epochs, etc. Beyond this we also test the effects of data augmentation.

2.1 ShallowNet

For baseline comparisons we will use the code provided by Dr. Rosebrock [1] which is a simplistic implementation of a Convolution Neural Network designed with the goal of classification of images in to any of the three given animal categories. The basic structure of Rosebrock’s network consists of:

INPUT=>CONV=>RELU=>FC=>SOFTMAX

Even with the most basic of implementations, the network is able to provide a respectable first attempt at the classification of the images. As seen in Table 1, the network achieves an average accuracy 66%. Given that with three classes the network has a 33% probability of randomly selecting the correct class, this is a decent first attempt. Using this as a lower bound we proceed to implement a number of other techniques that are common in CNN architectures in order to improve the network performance overall.




Figure 1: ShallowNet Accuracy

	Precision	Recall	F1	Support
Cat	0.64	0.57	0.6	262
Dog	0.57	0.58	0.58	249
Panda	0.79	0.86	0.82	239
Avg/Tot	0.66	0.67	0.66	750


Table 1: ShallowNet Results

2.2 MohlerNet2

The initial modification to the original network lies in the expansion of the architecture to incorporate key elements of CNNs, such as pooling layers and incorporation of neuron dropout prior to the fully connected layer. The general architecture of “MohlerNet2” is as follows:

INPUT=>CONV=>CONV=>MAXPOOL=>DROPOUT(0.5)=>FC=>SOFTMAX

More specifically, the CNN employs the use of 3×3 kernels in the convolution layers, where the first layer generates 32 feature maps and the second expands to 64 feature maps prior to applying a 2×2 max-pooling to the output. From the pooled output a neuron dropout probability of 50% is applied in order to attempt to reduce neuron dependencies and memorization of data. Using the classical stochastic gradient descent (SGD) approach to the optimization of the network we were able to see an immediate and reproducible increase in the ability of the network to accurately classify the three categories of animal images. From Table 2, it can be seen that the inclusion of the additional convolutional layer, pooling, and dropout was able to provide a 6% increase in the accuracy of the network. For this experiment we used the keras recommended learning rate ($LR = 0.01$). As a more objective measure of the algorithm we can view the F-1 score for each tested network. The F-1 score captures information regarding the false positives and false negatives in the classification process, this gives a more objective view from algorithm to algorithm. From this we can also see a similar increase in F-1 score between the original ShallowNet implementation and the first implementation of MohlerNet (MohlerNet2).




	Precision	Recall	F1	Support
Cat	0.67	0.75	0.71	262
Dog	0.64	0.60	0.62	249
Panda	0.88	0.82	0.85	239
Avg/Tot	0.72	0.72	0.72	750

Table 2: MohlerNet2 Classification Results

Figure 2: MohlerNet2 Accuracy, Learning Rate:0.01

Once the first trials of the SGD optimized results were obtained we observed the effect of modifying the optimizer for the network. While using the AdaGrad and Adam optimizers, which apply adaptive learning rates, we saw no improvement in any of the scoring metrics. Actually, there was a reduction in accuracy from the original SGD approach, down to an average of approximately 70% classification accuracy. In order to attempt to increase the

accuracy of the network we must take in to account that we are dealing with a relatively small data set (by modern standards). With the small number of images being used to train the network there is a high probability that we will experience over-fitting while attempting to increase performance. In order to counter this, we next applied data augmentation to the training set. This allows us to modify the images in moderate ways (rotations, flips, sheering, etc.) such that they are presented to the network in a new way each epoch. The inclusion of data augmentation yielded another (relatively) considerable increase in accuracy and F1-Score. As shown in Table 3, both metrics increased by 4% such that the MohlerNet2 implementation could classify with 76% accuracy. However, from Figure 3 it can be seen that from the gap between training accuracy and validation accuracy, the network is experiencing some mild over fitting, which indicates that some modifications may be able to improve the performance further, these remain to be discovered.



	Precision	Recall	F1	Support
Cat	0.75	0.66	0.71	262
Dog	0.69	0.7	0.69	249
Panda	0.84	0.93	0.89	239
Avg/Tot	0.76	0.76	0.76	750

Table 3: MohlerNet2 with Data Augmentation Results

Figure 3: MohlerNet2 with Data Augmentation Accuracy, AdaGrad Optimized

2.3 MohlerNet3

The next generation implementation of MohlerNet (MohlerNet3) is based upon the architecture of Alex-Net [2]

INPUT=>[CONV=>MAXPOOL]*2=>CONV*3=>MAXPOOL=>DROPOUT(0.5)=>FC=>SOFTMAX

2.4 MohlerNet4

INPUT=>[CONV=>MAXPOOL]*3=>FC=>DROPOUT=>FC=>SOFTMAX

3 Conclusions




Figure 4: MohlerNet3 with Data Augmentation Accuracy, AdaMax Optimized

	Precision	Recall	F1	Support
Cat	0.81	0.58	0.68	262
Dog	0.62	0.78	0.69	249
Panda	0.88	0.91	0.89	239
Avg/Tot	0.77	0.75	0.75	750

Table 4: MohlerNet3 with Data Augmentation Results




Figure 5: MohlerNet3 with Data Augmentation Accuracy, AdaMax Optimized

	Precision	Recall	F1	Support
Cat	0.81	0.58	0.68	262
Dog	0.62	0.78	0.69	249
Panda	0.88	0.91	0.89	239
Avg/Tot	0.77	0.75	0.75	750

Table 5: MohlerNet3 with Data Augmentation Results

A Code Listings

Listing 1: MohlerNet Implementations

```

1 # import the necessary packages
  from keras.models import Sequential
3 from keras.layers.convolutional import Conv2D,MaxPooling2D
  from keras.layers.core import Activation
5 from keras.layers.core import Flatten,Dropout
  from keras.layers.core import Dense
7 from keras import regularizers
  from keras import backend as K
9

```

```

class MohlerNet1:
11     @staticmethod
12     def build(width, height, depth, classes):
13         # initialize the model along with the input shape to be
14         # "channels last"
15         model = Sequential()
16         inputShape = (height, width, depth)
17
18         # if we are using "channels first", update the input shape
19         if K.image_data_format() == "channels_first":
20             inputShape = (depth, height, width)
21
22         # define the first (and only) CONV => RELU layer
23         model.add(Conv2D(32, (3, 3), padding="same",
24                         input_shape=inputShape))
25         model.add(Activation("relu"))
26
27         # softmax classifier
28         model.add(Flatten())
29         model.add(Dense(classes))
30         model.add(Activation("softmax"))
31
32         # return the constructed network architecture
33         return model
34
35 #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
class MohlerNet2:
37     @staticmethod
38     def build(width, height, depth, classes):
39         # initialize the model along with the input shape to be
40         # "channels last"
41         model = Sequential()
42         inputShape = (height, width, depth)
43
44         # if we are using "channels first", update the input shape
45         if K.image_data_format() == "channels_first":
46             inputShape = (depth, height, width)
47
48         # define the first (and only) CONV => RELU layer
49         model.add(Conv2D(32, (3, 3), padding="same", activation='
50             [+]relu',
51             input_shape=inputShape))
52         #model.add(Activation("relu"))
53
54         #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT

```

```

        model.add(Conv2D(64,(3,3),padding="same",activation='relu'
            [+]))
55     model.add(MaxPooling2D(pool_size=(2,2)))
        model.add(Dropout(0.5))
57
        regularizers.l2
59     # softmax classifier
        model.add(Flatten())
61     model.add(Dense(classes))
        model.add(Activation("softmax"))
63
        # return the constructed network architecture
65     return model

67 class MohlerNet3:
    @staticmethod
69     def build(width, height, depth, classes):
        # initialize the model along with the input shape to be
71         # "channels last"
        model = Sequential()
73         inputShape = (height, width, depth)

75         # if we are using "channels first", update the input shape
        if K.image_data_format() == "channels_first":
77             inputShape = (depth, height, width)

79         # define the first (and only) CONV => RELU layer
        model.add(Conv2D(64, (3, 3), padding="same", activation='
            [+]'relu',
81             input_shape=inputShape))
        model.add(MaxPooling2D(pool_size=(2,2)))
83
        #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
85     model.add(Conv2D(128,(3,3),padding="same",activation='relu'
            [+]))
        model.add(MaxPooling2D(pool_size=(2,2)))
87
        model.add(Conv2D(64,(3,3),padding="same",activation='relu'
            [+]))
89     model.add(Conv2D(32,(3,3),padding="same",activation='relu'
            [+]))
        model.add(Conv2D(32,(3,3),padding="same",activation='relu'
            [+]))
91     model.add(MaxPooling2D(pool_size=(2,2)))

```

```

93     model.add(Flatten())
94     model.add(Dense(classes))
95     model.add(Dropout(0.5))
96     model.add(Activation("softmax"))
97
98     # return the constructed network architecture
99     return model
100
101
102 class MohlerNet4:
103     @staticmethod
104     def build(width, height, depth, classes):
105         # initialize the model along with the input shape to be
106         # "channels last"
107         model = Sequential()
108         inputShape = (height, width, depth)
109
110         # if we are using "channels first", update the input shape
111         if K.image_data_format() == "channels_first":
112             inputShape = (depth, height, width)
113
114         # define the first (and only) CONV => RELU layer
115         model.add(Conv2D(32, (3, 3), padding="same", activation='
            [+]relu',
            input_shape=inputShape))
116         model.add(MaxPooling2D(pool_size=(2,2)))
117         #model.add(Activation("relu"))
118
119         #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
120         model.add(Conv2D(32,(3,3),padding="same",activation='relu'
            [+]))
121         model.add(MaxPooling2D(pool_size=(2,2)))
122
123         model.add(Conv2D(64,(3,3),padding="same",activation='relu'
            [+]))
124         model.add(MaxPooling2D(pool_size=(2,2)))
125         # softmax classifier
126         model.add(Flatten())
127         model.add(Dense(64))
128         model.add(Activation("relu"))
129         model.add(Dropout(0.5))
130         model.add(Dense(classes))
131         model.add(Activation("softmax"))
132
133         # return the constructed network architecture

```


References

- [1] A. Rosebrock, *Deep Learning For Computer Vision With Python*. PyImageSearch, 2017.
- [2] A. Krizhevsky, I. Sutskever, and G. E. Hinton, “Imagenet classification with deep convolutional neural networks,” in *Advances in Neural Information Processing Systems 25* (F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, eds.), pp. 1097–1105, Curran Associates, Inc., 2012.