

# EE 5450 Project 02: Convolution Neural Network Classification of the Animals Data Set

David R. Mohler

April 18, 2018

## 1 Introduction

The rebirth of Neural Networks (NNs) and Machine Learning (ML) has offered revolutionary approaches to intelligent tasks previously unattainable in the field of image processing and computer vision. A major area of research continues to be the automated classification of images. Image classification through classical means of image processing tend to suffer due to problems with the wide variance in images to be classified and how their subject matter may be transformed from image to image. The discovery of Convolution Neural Networks (CNNs) lead to a massive increase in the accuracy of which images could be classified in an automated fashion.

For this project we will demonstrate the development of several convolution neural networks that vary across a wide number of parameters for the purpose of classifying color images of three different types of animals. The CNN is tasked with classifying images of cats, dogs, and pandas, typical examples of images in the data set can be seen below. Through the variation of network parameters such as the number of layers and the number of neurons within them, activation functions, learning rates, optimization routines, data augmentation, and other hyper parameters, we display the ability to increase the classification performance of the network.



## 2 Methods and Results

After establishing an initial understanding of the basic **ShallowNet** animal classification network that was provided, we begin with a relatively simple and small expansion to include fundamental elements of CNNs, once a working network is establish we are able to tune the parameters and observe the effects of varying the optimization techniques. Once a candidate optimizer has been selected we proceed to modify and tune the associated hyperparameters such as the learning rate (when applicable), batch size, number of training epochs, etc. Beyond this we also test the effects of data augmentation.

### 2.1 ShallowNet

For baseline comparisons we will use the code provided by Dr. Rosebrock [1] which is a simplistic implementation of a Convolution Neural Network designed with the goal of classification of images in to any of the three given animal categories. The basic structure of Rosebrock’s network consists of:

INPUT=>CONV=>RELU=>FC=>SOFTMAX

Even with the most basic of implementations, the network is able to provide a respectable first attempt at the classification of the images. As seen in Table 1, the network achieves an average accuracy 66%. Given that with three classes the network has a 33% probability of randomly selecting the correct class, this is a decent first attempt. Using this as a lower bound we proceed to implement a number of other techniques that are common in CNN architectures in order to improve the network performance overall.




Figure 1: ShallowNet Accuracy

	Precision	Recall	F1-Score	Support
Cat	0.64	0.57	0.6	262
Dog	0.57	0.58	0.58	249
Panda	0.79	0.86	0.82	239
Avg/Total	0.66	0.67	0.66	750


Table 1: ShallowNet Results

## 2.2 MohlerNet2

The initial modification to the original network lies in the expansion of the architecture to incorporate key elements of CNNs, such as pooling layers and incorporation of neuron dropout prior to the fully connected layer. The general architecture of “MohlerNet2” is as follows:

INPUT=>CONV=>CONV=>MAXPOOL=>DROPOUT(0.5)>FC=>SOFTMAX

Using the classical stochastic gradient descent (SGD) approach to the optimization of the network we were able to see an immediate and reproducible increase in the ability of the network to accurately classify the three categories of animal images. From Table 2, it can be seen that the inclusion of the additional convolutional layer, pooling, and dropout was able to provide a 6% increase in the accuracy of the network. For this experiment we used the keras recommended learning rate ( $LR = 0.01$ ). As a more objective measure of the algorithm we can view the F-1 score for each tested network. The F-1 score captures information regarding the false positives and false negatives in the classification process, this gives a more objective view from algorithm to algorithm. From this we can also see a similar increase in F-1 score between the original ShallowNet implementation and the first implementation of MohlerNet (MohlerNet2).



	Precision	Recall	F1-Score	Support
Cat	0.67	0.75	0.71	262
Dog	0.64	0.60	0.62	249
Panda	0.88	0.82	0.85	239
Avg/Total	0.72	0.72	0.72	750

Table 2: MohlerNet2 Classification Results

Figure 2: MohlerNet2 Accuracy, Learning Rate:0.01

Once the first results of the SGD optimized results were obtained we observed the effect of modifying the optimizer for the network. While using the AdaGrad and Adam optimizers, which apply adaptive learning rates, we saw no improvement in any of the scoring metrics. Actually, there was a reduction in accuracy from the original SGD approach, down to an average of approximately 70% classification accuracy.

**TALK ABOUT OVERFITTING HERE**

## 2.3 MohlerNet3

INPUT=> [CONV=>MAXPOOL] \*2=>CONV\*3=>MAXPOOL=>DROPOUT(0.5)>FC=>SOFTMAX




Figure 3: MohlerNet2 with Data Augmentation Accuracy, AdaGrad Optimized

	Precision	Recall	F1-Score	Support
Cat	0.75	0.66	0.71	262
Dog	0.69	0.7	0.69	249
Panda	0.84	0.93	0.89	239
Avg/Total	0.76	0.76	0.76	750

Table 3: MohlerNet2 with Data Augmentation Results




Figure 4: MohlerNet3 with Data Augmentation Accuracy, AdaMax Optimized


	Precision	Recall	F1-Score	Support
Cat	0.81	0.58	0.68	262
Dog	0.62	0.78	0.69	249
Panda	0.88	0.91	0.89	239
Avg/Total	0.77	0.75	0.75	750

Table 4: MohlerNet3 with Data Augmentation Results

## 2.4 MohlerNet4

INPUT=> [CONV=>MAXPOOL] \*3=>FC=>DROPOUT=>FC=>SOFTMAX

## 3 Conclusions



	Precision	Recall	F1-Score	Support
Cat	0.81	0.58	0.68	262
Dog	0.62	0.78	0.69	249
Panda	0.88	0.91	0.89	239
Avg/Total	0.77	0.75	0.75	750

Table 5: MohlerNet3 with Data Augmentation Results

Figure 5: MohlerNet3 with Data Augmentation Accuracy, AdaMax Optimized

## A Code Listings

Listing 1: MohlerNet Implementations

```

1 # import the necessary packages
  from keras.models import Sequential
3 from keras.layers.convolutional import Conv2D,MaxPooling2D
  from keras.layers.core import Activation
5 from keras.layers.core import Flatten,Dropout
  from keras.layers.core import Dense
7 from keras import regularizers
  from keras import backend as K
9
class MohlerNet1:
11     @staticmethod
    def build(width, height, depth, classes):
13         # initialize the model along with the input shape to be
        # "channels last"
15         model = Sequential()
        inputShape = (height, width, depth)
17
        # if we are using "channels first", update the input shape
19         if K.image_data_format() == "channels_first":
            inputShape = (depth, height, width)
21
        # define the first (and only) CONV => RELU layer
23         model.add(Conv2D(32, (3, 3), padding="same",
            input_shape=inputShape))
25         model.add(Activation("relu"))

```

```

27     # softmax classifier
    model.add(Flatten())
29     model.add(Dense(classes))
    model.add(Activation("softmax"))
31
    # return the constructed network architecture
33     return model

35 #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
class MohlerNet2:
37     @staticmethod
    def build(width, height, depth, classes):
39         # initialize the model along with the input shape to be
        # "channels last"
41         model = Sequential()
        inputShape = (height, width, depth)
43
        # if we are using "channels first", update the input shape
45         if K.image_data_format() == "channels_first":
            inputShape = (depth, height, width)
47
        # define the first (and only) CONV => RELU layer
49         model.add(Conv2D(32, (3, 3), padding="same", activation='
            [+]relu',
                input_shape=inputShape))
51         #model.add(Activation("relu"))

53         #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
        model.add(Conv2D(64, (3, 3), padding="same", activation='relu'
            [+]))
55         model.add(MaxPooling2D(pool_size=(2, 2)))
        model.add(Dropout(0.5))
57
        regularizers.l2
59         # softmax classifier
        model.add(Flatten())
61         model.add(Dense(classes))
        model.add(Activation("softmax"))
63
        # return the constructed network architecture
65         return model

67 class MohlerNet3:
    @staticmethod

```

```

69     def build(width, height, depth, classes):
70         # initialize the model along with the input shape to be
71         # "channels last"
72         model = Sequential()
73         inputShape = (height, width, depth)
74
75         # if we are using "channels first", update the input shape
76         if K.image_data_format() == "channels_first":
77             inputShape = (depth, height, width)
78
79         # define the first (and only) CONV => RELU layer
80         model.add(Conv2D(64, (3, 3), padding="same", activation='
81             [+]relu',
82             input_shape=inputShape))
83         model.add(MaxPooling2D(pool_size=(2,2)))
84
85         #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
86         model.add(Conv2D(128,(3,3),padding="same",activation='relu
87             [+]'))
88         model.add(MaxPooling2D(pool_size=(2,2)))
89
90         model.add(Conv2D(64,(3,3),padding="same",activation='relu'
91             [+]))
92         model.add(Conv2D(32,(3,3),padding="same",activation='relu'
93             [+]))
94         model.add(Conv2D(32,(3,3),padding="same",activation='relu'
95             [+]))
96         model.add(MaxPooling2D(pool_size=(2,2)))
97
98         model.add(Flatten())
99         model.add(Dense(classes))
100         model.add(Dropout(0.5))
101         model.add(Activation("softmax"))
102
103         # return the constructed network architecture
104         return model
105
106 class MohlerNet4:
107     @staticmethod
108     def build(width, height, depth, classes):
109         # initialize the model along with the input shape to be
110         # "channels last"
111         model = Sequential()
112         inputShape = (height, width, depth)

```

```

109     # if we are using "channels first", update the input shape
111     if K.image_data_format() == "channels_first":
112         inputShape = (depth, height, width)
113
114     # define the first (and only) CONV => RELU layer
115     model.add(Conv2D(32, (3, 3), padding="same", activation='
        [+]relu',
        input_shape=inputShape))
117     model.add(MaxPooling2D(pool_size=(2,2)))
118     #model.add(Activation("relu"))
119
120     #ADD CONVOLUTION LAYER, POOLING LAYER, AND DROPOUT
121     model.add(Conv2D(32,(3,3),padding="same",activation='relu'
        [+]))
122     model.add(MaxPooling2D(pool_size=(2,2)))
123
124     model.add(Conv2D(64,(3,3),padding="same",activation='relu'
        [+]))
125     model.add(MaxPooling2D(pool_size=(2,2)))
126     # softmax classifier
127     model.add(Flatten())
128     model.add(Dense(64))
129     model.add(Activation("relu"))
130     model.add(Dropout(0.5))
131     model.add(Dense(classes))
132     model.add(Activation("softmax"))
133
134     # return the constructed network architecture
135     return model

```



## References

- [1] A. Rosebrock, *Deep Learning For Computer Vision With Python*. PyImageSearch, 2017.